

How the Proof Works

The three-dimensional Navier–Stokes regularity problem asks whether a smooth flow can become singular in finite time. We begin with a smooth, incompressible flow of finite energy and a fixed positive viscosity. Its velocity and pressure satisfy

$$\underbrace{\partial_t u}_{\text{local acceleration}} + \underbrace{(u \cdot \nabla) u}_{\text{nonlinear transport}} = \underbrace{-\nabla p}_{\text{pressure}} + \underbrace{\nu \Delta u}_{\text{viscous diffusion}}, \quad \nabla \cdot u = 0.$$

The initial velocity is $u_0 \in H^\infty_\sigma(\mathbb{R}^3)$, and $\nu > 0$ stays fixed throughout. A classical solution exists for a while; the question is whether its maximal lifespan $[0, T_*)$ can end at a finite time. The hope is simple: viscosity smooths the fluid faster than nonlinear transport can sharpen it.

Strain within the fluid’s flow can naturally stretch a vortex. When that happens, incompressibility forces a material tube to narrow as it lengthens. If circulation persists strongly enough through that narrowing, the fluid circulates faster around a smaller core. The velocity then changes more sharply across the core, and the stretching can intensify the very gradients that viscosity is smoothing. Viscosity diffuses momentum between neighboring regions of fluid, reducing velocity differences. Whether it can smooth those differences fast enough is the difficulty.

A narrowing vortex makes that difficulty progressively more local. Large gradients can occupy a region too small to account for much of the total energy, while both nonlinear transport and viscous diffusion become stronger there. Following the concentration toward vanishing width forces us to compare their behavior at smaller and smaller scales. We need to know whether viscosity gains on transport as the width tends to zero.

The equations allow an exact comparison. For a scale factor $\lambda > 1$, shrink lengths by λ , increase velocities by λ , and shorten times by λ^2 :

$$u_\lambda(x, t) = \lambda u(\lambda x, \lambda^2 t), \quad p_\lambda(x, t) = \lambda^2 p(\lambda x, \lambda^2 t).$$

The time derivative, nonlinear transport, pressure gradient, and viscous term all acquire a factor λ^3 . After that common factor is removed, the viscosity is still ν . The narrower, faster flow obeys the original equation with precisely the original balance of scaling weights. Shrinking the scale has supplied no extra viscous advantage. Which effect dominates the evolution still depends on the fluid’s actual configuration.

To measure what this concentration does to spatial regularity, the Sobolev index s lets us vary how strongly we count fine spatial variation. At integer values it counts square-integrable derivatives; fractional values interpolate that sensitivity by weighting a spatial frequency ξ by $|\xi|^s$. The velocity factor contributes λ , the derivative weight contributes λ^s , and the square root of the volume factor contributes $\lambda^{-3/2}$. Together they give

$$\|u_\lambda(\cdot, t)\|_{\dot{H}^s} = \lambda^{s-1/2} \|u(\cdot, \lambda^2 t)\|_{\dot{H}^s}.$$

Below $s = \frac{1}{2}$, concentration makes this norm smaller. Above $s = \frac{1}{2}$, it makes the norm larger. Exactly at $s = \frac{1}{2}$, the increased sensitivity to fine variation offsets the shrinking volume, and the norm is unchanged. This is the critical height. Its invariance compares rescaled flows; it is not a conservation law along an evolving flow.

Energy lies below that height. It is half the squared L^2 norm, so it scales as λ^{-1} : the velocity grows, but the volume containing the fast motion shrinks even faster. Increasing concentration and speed can coexist with falling total energy. The energy identity can keep the total finite without, by that fact alone, excluding an arbitrarily narrow and fast region.

The time available for further concentration also shrinks. At width ℓ and velocity scale U , transport acts over a time of order ℓ/U ; under this rescaling that time gains a factor λ^{-2} . The viscous time ℓ^2/ν gains that factor too. If the evolution could repeat such stages, each narrower stage would take less time than the last:

$$\ell_j = \lambda^{-j} \ell_0, \quad U_j = \lambda^j U_0, \quad \tau_j = \lambda^{-2j} \tau_0, \quad \sum_{j=0}^{\infty} \tau_j = \frac{\tau_0}{1 - \lambda^{-2}} < \infty.$$

Infinitely many stages could then fit before a finite terminal time. That is the hypothetical Zeno cascade: narrowing without a smallest width, acceleration without a final stage, and finite energy throughout. Scaling allows this comparison. It does not establish that a Navier–Stokes flow actually performs the cascade.

For a finite-time singularity, velocity would have to blow up. That requires acceleration to blow up, then jerk, and so on, and so on. Here we follow the velocity field at fixed spatial positions, so the sequence is $u, \partial_t u, \partial_t^2 u, \dots$. These are Eulerian time derivatives; a moving particle has material acceleration $D_t u = \partial_t u + (u \cdot \nabla)u$, and its later responses use $D_t = \partial_t + u \cdot \nabla$. The necessity has a precise meaning in the spatial L^∞ norm. For any fixed $t_0 < T_*$ and every finite order j , integration gives

$$\|\partial_t^j u(t)\|_\infty \leq \|\partial_t^j u(t_0)\|_\infty + \int_{t_0}^t \|\partial_t^{j+1} u(s)\|_\infty ds.$$

A bounded next derivative could accumulate only a finite change before T_* . Unbounded velocity requires an infinite accumulation of its time derivative, which requires that derivative to be unbounded; applying the argument again forces the next order to be unbounded as well. The norms need not grow monotonically, and their large values need not occur at one fixed point. The requirement persists through every finite temporal order.

But these increasingly demanding time derivatives are still derivatives of a velocity governed by Navier–Stokes. Each time we differentiate, the viscous term applies a Laplacian to the preceding response. The first time derivative contains Δu , the second contains $\Delta \partial_t u$, and substitution of the first equation brings in $\Delta^2 u$. Continuing exposes higher spatial derivatives, together with all the products created by differentiating transport and pressure. The temporal escalation has reached into the spatial structure of the flow.

We already measured the velocity’s fine spatial variation with the critical norm. Differentiating that norm follows how this spatial variation changes in time, so it brings the temporal responses directly into the measurement. With $\Lambda = (-\Delta)^{1/2}$, write half its squared size as $H = \frac{1}{2} \|\Lambda^{1/2} u\|_2^2$. Differentiating the two velocity factors gives

$$\begin{aligned} H' &= \text{Re} \langle \Lambda^{1/2} u, \Lambda^{1/2} \partial_t u \rangle, \\ H'' &= \|\Lambda^{1/2} \partial_t u\|_2^2 + \text{Re} \langle \Lambda^{1/2} u, \Lambda^{1/2} \partial_t^2 u \rangle. \end{aligned}$$

Acceleration and jerk enter through these actual differentiated velocity factors. The equation now expresses those factors through spatial derivatives, transport, and pressure. In the first pairing, its viscous contribution is

$$\nu \langle \Lambda^{1/2} u, \Lambda^{1/2} \Delta u \rangle = -\nu \|\Lambda^{3/2} u\|_2^2.$$

The Laplacian has converted the temporal response into a measurement of finer spatial variation. Integration by parts does this at any Sobolev height: if $E_s = \frac{1}{2} \|\Lambda^s u\|_2^2$, then

$$E'_s = -2\nu E_{s+1} + \mathcal{P}_s, \quad \mathcal{P}_s := -\text{Re} \langle \Lambda^s u, \Lambda^s ((u \cdot \nabla)u + \nabla p) \rangle_{L^2}.$$

Differentiate the critical-height relation again. The new spatial energy $E_{3/2}$ changes according to this identity too, exposing $E_{5/2}$; its transport and pressure terms also have to be differentiated. The first two orders become

$$\begin{aligned} H' &= -2\nu E_{3/2} + \mathcal{P}_{1/2}, \\ H'' &= 4\nu^2 E_{5/2} - 2\nu \mathcal{P}_{3/2} + \partial_t \mathcal{P}_{1/2}. \end{aligned}$$

The two expressions for H'' concern one quantity: the pairing of velocity with its higher temporal responses also contains its higher spatial Sobolev energy, with the differentiated nonlinear contribution retained. Continuing repeats this relation at every finite order. The temporal responses and spatial regularity cannot be specified independently of one another; the equation couples them through these identities. Their bounds must account for the nonlinear terms as well as the viscous term.

Those spatial demands are exactly what smoothness needs. A bound at one Sobolev height controls only a finite amount of spatial variation. Raise the height far enough, however, and Sobolev embedding makes any prescribed spatial derivative continuous: to obtain k continuous derivatives, choose $m > k + \frac{3}{2}$. If the velocity belongs to every finite height, we can make that choice for every k :

$$L^2 \cap \bigcap_{n \geq 0} \dot{H}^{n+1/2} = \bigcap_{m \geq 0} H^m = H^\infty \hookrightarrow C^\infty.$$

No single finite height gives all of smoothness. The intersection does, because every derivative we ask for has a sufficiently high finite space in which it is controlled. The constants may increase with the order; smoothness asks that each finite order exists, with its own finite bound.

The equation also explains why this spatial intersection can accommodate the entire temporal sequence. A Laplacian uses two spatial derivatives; at any requested height, the intersection has those additional derivatives available. Products of smooth Sobolev fields and the pressure recovered from them also remain in the intersection. Once the spatial and temporal requirements have been realized together over the approach to T_* , the equation determines the next temporal response there, and its differentiated equations continue that relation through every finite order. The time and space requirements meet in a velocity with enough spatial regularity for every temporal response and enough time regularity to reach a common endpoint. That endpoint is smooth and can serve as initial data for continued evolution. This is the connection the proof has to realize from u_0 ; the differentiated identities reveal it, while the construction must establish the bounds and compatibility that make it true at the endpoint.

If a singularity were to form, we would expect the opposite behavior on the approach to that time. The velocity could remain smooth at every earlier instant while its increasingly fine structure forced some fixed spatial Sobolev norm beyond every bound. Its temporal responses would become unbounded as well, even while the total energy remained finite. We would have all finite derivatives before T_* , yet lose the control needed to carry them through T_* . The distinction is between smoothness on each shorter interval and regularity maintained over the full approach to the putative singularity. A proof has to exclude this loss at the endpoint.

To do that, we return to the derivatives the initial velocity actually generates. Pressure is fixed by incompressibility and its decay normalization: taking the divergence of the equation gives $-\Delta p = \partial_i \partial_j (u_i u_j)$. Differentiating the velocity equation now gives the first three responses in their uncompressed form:

$$\partial_t u = \nu \Delta u - (u \cdot \nabla) u - \nabla p,$$

$$\partial_t^2 u = \nu \Delta \partial_t u - (\partial_t u \cdot \nabla) u - (u \cdot \nabla) \partial_t u - \nabla \partial_t p,$$

$$\partial_t^3 u = \nu \Delta \partial_t^2 u - (\partial_t^2 u \cdot \nabla) u - 2(\partial_t u \cdot \nabla) \partial_t u - (u \cdot \nabla) \partial_t^2 u - \nabla \partial_t^2 p.$$

The factor of two comes from differentiating both factors of transport. At the next order each product splits again. Writing $V_r = \partial_t^r u$ and $Q_r = \partial_t^r p$ lets us retain all those terms without expanding an ever longer line:

$$V_{r+1} = \nu \Delta V_r - \sum_{a=0}^r \binom{r}{a} (V_a \cdot \nabla) V_{r-a} - \nabla Q_r,$$

$$Q_r = R_i R_j \sum_{a=0}^r \binom{r}{a} V_{a,i} V_{r-a,j}.$$

Here R_i are the Riesz transforms and repeated spatial indices are summed. The pressure formula differentiates the incompressibility constraint along with the velocity. At $t = 0$, these relations start with $V_0 = u_0$ and determine each subsequent velocity and pressure derivative. Spatial differentiation fills out each temporal row. This collection is the mixed jet: the time and space derivatives generated by the datum under the equation.

We need only finitely many of those derivatives for any specified order of regularity. Fix a temporal depth N and a spatial height M . The datum-generated rectangle theorem, Theorem 4.6, supplies the adjacent rows on the whole interval ending at T_* . In the cutoff form of its conclusion,

$$\sup_{0 < T < T_*} \sum_{n=0}^N \left(\|V_n\|_{L^2(0,T;H^{M+1})}^2 + \|V_{n+1}\|_{L^2(0,T;H^{M-1})}^2 \right) < \infty.$$

The bound can depend on N, M, u_0, ν , and the putative finite endpoint, but it holds as the cutoff T approaches that endpoint. Increasing the derivative orders may increase the bound; approaching T_* at any fixed pair of orders cannot make it diverge. The differentiated recurrence identifies every entry with a derivative of the original history, while the datum theorem supplies its full-interval control.

Letting T increase to T_* gives $V_n \in L^2(0, T_*; H^{M+1})$ and $\partial_t V_n = V_{n+1} \in L^2(0, T_*; H^{M-1})$. These adjacent spaces are chosen for a reason. The time derivative can be paired with V_n across the intermediate space H^M , and the two square-integrable factors give an integrable change in its squared norm. The Hilbert-triple trace theorem makes this precise:

$$H^{M+1} \subset H^M \subset H^{M-1}, \quad V_n \in C([0, T_*]; H^M).$$

Each retained velocity derivative now has an actual endpoint value at the requested spatial height. The pressure formula supplies the corresponding pressure rows. We collect these finitely many velocity, adjacent time, and pressure rows into a rectangle $\mathcal{R}_{N,M}[u_0; T_*]$. Its two indices remain independent, since a request for another time derivative and a request for finer spatial regularity are different requests.

Increasing either depth retains every derivative already present. An entry in two rectangles is a distributional derivative of the original velocity or pressure in both, so their restrictions agree. This matters at the endpoint: different spatial heights must describe one limiting field, and different temporal depths must describe its actual derivatives. Compatibility lets all the finite rectangles be retained together:

$$\mathcal{I}[u_0] := (\mathcal{R}_{N,M}[u_0; T_*])_{N,M \geq 0}, \quad u \in \bigcap_{r,m \geq 0} H^r(0, T_*; H^m(\mathbb{R}^3)).$$

For any finite r, m , choose a rectangle deep enough to contain those derivatives. Every larger rectangle preserves them. The projective intersection expresses this agreement through all finite orders; it does not require summing their bounds into one infinite norm.

We can now take the endpoint in every spatial height. The zeroth temporal row gives $u \in C([0, T_*]; H^m)$ for every finite m , and compatibility identifies its traces as one field:

$$u_* := \lim_{t \uparrow T_*} u(t) \quad \text{in every } H^m, \quad u_* \in H_\sigma^\infty \hookrightarrow C^\infty.$$

The differentiated velocity and pressure relations also reach this endpoint, determining its full jet from u_* . Smooth local existence then starts a solution from u_* . Local uniqueness on overlapping time intervals joins that solution to the original flow approaching T_* , continuing it beyond its supposed maximal time. A finite T_* is impossible, so $T_* = \infty$.

With the intersection constructed, we can derive the finite estimates, and use them. Any estimate requiring finitely many derivatives can take its regularity from a sufficiently large rectangle. For a finite time T , temporal order r , spatial order k , and $2 \leq p \leq \infty$, choose $m > k + \frac{3}{2} - \frac{3}{p}$. The corresponding coordinate and Sobolev embedding give

$$\partial_t^r u \in C([0, T]; H^m) \hookrightarrow L^\infty(0, T; W^{k,p}).$$

Taking $r = 0$, $k = 1$, $p = \infty$, and $m = 3$, for example, gives

$$\int_0^T \|\nabla u(t)\|_\infty dt \leq T \sup_{0 \leq t \leq T} \|\nabla u(t)\|_\infty < \infty.$$

The accumulated velocity gradient is finite on every finite interval. This is the familiar control that a concentrating vortex would have to defeat, now obtained from a finite projection of the completed intersection.